

OpenStreetMap - Data Wrangling

WGU | Data Wrangling

Udacity | Project: OpenStreetMap Data

Full Project Document

The abbreviated project document without code blocks, output, or tables can be found here:

[osmAustin project docs](#)

Purpose

This project was created for Udacity's Data Analyst Nanodegree. An extract of xml data was downloaded for a selected city or region from OpenStreetMap (OSM). This document details the auditing, cleaning, transformation, and analysis performed on that raw dataset. After the raw data was cleaned and staged in a tabular format (csv), it was loaded into a database for additional analysis.

Selecting a Dataset

For this project I decided to work with data from Austin, TX. The selected map area is too large to export directly from OpenStreetMap,¹ but I found a suitable extract hosted by Interline.² This particular extract was a pbf file, so I had to convert it to osm file format before auditing the data. I used a command line tool called osmosis to make this conversion.³

```
# install homebrew
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)

# install osmosis
brew install osmosis

# convert pbf to osm
osmosis --read-pbf \austin_texas.osm.pbf --write-xml austin_texas.osm
```

Auditing

To begin the auditing process, I created three summaries; one each for elements, attributes, and keys. I created four functions to accomplish this: `print_sorted_dict`, `count_elements`, `count_attributes`, and `count_keys`.⁴

```

import xml.etree.ElementTree as ET

def print_sorted_dict(d, sort_by=None):
    """Prints dictionary sorted by keys or items"""
    if sort_by is None:
        sort_by = 'items'
    if sort_by == 'keys':
        sorted_dict = sorted(d.keys(), key=lambda s: s.lower())
    elif sort_by == 'items':
        sorted_dict = dict(sorted(d.items(), key=lambda s: s[1], reverse=True))
    else:
        print("Invalid sort_by: please input 'keys' or 'items'\n")
        sorted_dict = d
    for k in sorted_dict:
        v = d[k]
        print(f'{k}: {v}')

def count_elements(filename):
    """Prints element tag name and count for each XML element."""
    d = {}
    for event, elem in ET.iterparse(filename, events=('start',)):
        if elem.tag not in d:
            d[elem.tag] = 1
        else:
            d[elem.tag] += 1
    print('\n----- Count all tags -----')
    print_sorted_dict(d)
    return

def count_attributes(filename):
    """Prints attribute name and count for each XML element."""
    d = {}
    for event, elem in ET.iterparse(filename, events=('start', 'end')):
        if event == 'end':
            for attr in elem.attrib:
                if attr not in d:
                    d[attr] = 1
                else:
                    d[attr] += 1
    print('\n----- Count all attributes -----')
    print_sorted_dict(d)

def count_keys(filename):
    """Prints key name and count for each XML element."""
    d = {}
    for event, elem in ET.iterparse(filename, events=('start', 'end')):
        if event == 'end':
            key = elem.attrib.get('k')
            if key:
                if key not in d:
                    d[key] = 1

```

```

    else:
        d[key] += 1
print('\n----- Count all keys -----')
print_sorted_dict(d)

```

Using these functions, I printed a few summaries.

```

from osmAudit import count_elements, count_attributes, count_keys

# define filename
atx_filename = 'data/austin_texas.osm'

# get count of elements
count_elements(atx_filename)

# get count of attributes
count_attributes(atx_filename)

# get count of keys
count_keys(atx_filename)

```

```

----- Count all tags -----
nd: 8835948
node: 7932057
tag: 2967844
way: 858496
member: 58198
relation: 4341
osm: 1

```

```

----- Count all attributes -----
ref: 8894146
version: 8794895
id: 8794894
timestamp: 8794894
uid: 8794894
user: 8794894
changeset: 8794894
lat: 7932057
lon: 7932057
k: 2967844
v: 2967844
type: 58198
role: 58198
generator: 1

```

```
----- Count all keys -----  
building: 622302  
height: 441107  
addr:street: 345406  
addr:housenumber: 344583  
highway: 216664  
addr:postcode: 98282  
name: 73274  
service: 52447  
access: 41496  
tiger:county: 37785  
tiger:cfcc: 37703  
surface: 34399  
tiger:name_base: 33396  
tiger:name_type: 30904  
tiger:reviewed: 25054  
oneway: 25010  
power: 23772  
barrier: 19658  
addr:city: 19394  
addr:state: 19314  
...
```

The basic components of the OpenStreetMap data model are tags, and the most important to this project are:

- node - describes points in space
- way - describes area boundaries and features
- relation - describe how other elements work together
- tag - describes the element to which they are attached, and contains two attributes: key (k) and value (v)

Exploring Key Values

Next, I checked the top 10 keys - based on frequency of occurrence - to see where there are opportunities for data cleaning. I also checked a few others that look interesting, and I created a function to facilitate this portion of the audit, `key_val_counter`.^{[4](#)}

```

import xml.etree.ElementTree as ET

def key_val_counter(filename, key_name):
    """Prints key name and count for each instance of key_name"""
    d = {}
    for event, elem in ET.iterparse(filename, events=('start', 'end')):
        if event == 'end':
            key = elem.attrib.get('k')
            if key == key_name:
                val = elem.attrib.get('v')
                if val not in d:
                    d[val] = 1
                else:
                    d[val] += 1
    print('\n----- Count of values for key: ' + key_name + ' -----')
    print_sorted_dict(d)

```

Because I was working in a python notebook, I couldn't just loop through the keys I was investigating. The printed data would get truncated well before all the keys' values were displayed. Instead, I decided to run each key in its own cell.

Most of the keys' values for Austin, TX OpenStreetMap data were already very clean. I suspect there are other students and hobbyists who have completed similar projects.

```

from osmAudit import key_val_counter

# define filename
atx_filename = 'data/austin_texas.osm'

# print key value counts
key_val_counter(atx_filename, 'height')
key_val_counter(atx_filename, 'addr:street')
key_val_counter(atx_filename, 'addr:housenumber')
key_val_counter(atx_filename, 'highway')
key_val_counter(atx_filename, 'name')
key_val_counter(atx_filename, 'service')
key_val_counter(atx_filename, 'tiger:county')

```

Problem Tags

Although most of the top key-values are clean, there are a few with opportunities for cleaning or filtering. I'll outline how these tags were cleaned/filtered in the next section.

building

```

from osmAudit import key_val_counter

# define filename
atx_filename = 'data/austin_texas.osm'

# key -> building
key_val_counter(atx_filename, 'building')

```

----- Count of values for key: building -----

yes: 584823

house: 20797

apartments: 4389

detached: 2685

carport: 1824

retail: 954

roof: 926

commercial: 758

school: 691

residential: 565

...

stadium seating: 4 ←

civic: 4

ruins: 4

container: 4

temple: 3

sports_centre: 3 ←

covered area: 2 ←

...

big state electric: 1 ←

tree_house: 1

undefined: 1

Bing: 1 ←

shelter: 1

gas_station: 1

transportation: 1

Learning_Center/_Day_Care: 1 ←

...

There are a few things that need to be cleaned-up in the values for the building key.

- There are spaces where there should be underscores. A simple string replace will correct those.
- A few other entries are incorrect or ambiguous; I'll correct those with a dictionary replace.

postcode

```

from osmAudit import key_val_counter

# define filename
atx_filename = 'data/austin_texas.osm'

# key -> addr:postcode
key_val_counter(atx_filename, 'addr:postcode')

```

```

----- Count of values for key: addr:postcode -----
78645: 10893
78734: 5627
78660: 4560
78653: 3553
78641: 3276
78669: 3190
78754: 2820
78704: 2559
78746: 2527
78723: 2290
...
78953: 3 ←
78644: 2 ←
78754;78753: 2 ←
78704-5639: 1 ←
78758-7008: 1 ←
...

```

These data are mostly clean, but there are some post codes included that are not actually in Austin.[5](#) I'll filter those out while cleaning and staging the data.

surface

```

from osmAudit import key_val_counter

# define filename
atx_filename = 'data/austin_texas.osm'

# key -> surface
key_val_counter(atx_filename, 'surface')

```

```

----- Count of values for key: surface -----
asphalt: 21169
paved: 5156
concrete: 3893
unpaved: 1407
concrete:plates: 558 ←

```

```
ground: 518
gravel: 452
dirt: 391
paving_stones: 250
fine_gravel: 181
...
cobblestone: 6
yes: 5 ←
con: 3 ←
mud: 2
...
CR_127: 1 ←
paving_stones:30: 1 ←
creekbed_(rock): 1 ←
concrete,_dirt: 1 ←
Large_unattached_stones_laid_in_the_creek: 1 ←
woodchips: 1
f: 1 ←
```

The values for the surface key need some cleaning. For some of them, I can figure out what the user intended - I can clean those with a dictionary. Some other values are less clear, and I'll remove those tags with a list.

city

```
from osmAudit import key_val_counter

# define filename
atx_filename = 'data/austin_texas.osm'

# key -> addr:city
key_val_counter(atx_filename, 'addr:city')
```

```
----- Count of values for key: addr:city -----
Austin: 12095
Cedar Park: 1985
Pflugerville: 1137
Round Rock: 1012
Georgetown: 713
Leander: 452
Elgin: 437
Hutto: 298
Bastrop: 280
Kyle: 181
...
```



```
Lost Pines: 2
AUSTIN: 2 ←
Pfluggerville: 2 ←
Wells Branch: 2 ←
Barton Creek: 1 ←
Ste 128, Austin: 1 ←
San Gabriel Village Boulevard: 1 ←
Dale: 1
manor: 1 ←
Pepe's Tacos: 1 ←
N Austin: 1 ←
Manchaca,: 1 ←
Austin;austin: 1 ←
kyle: 1 ←
Tampa: 1
McNeil: 1
Smithville: 1
wimberley: 1 ←
Wimberly: 1 ←
Marble Falls: 1
georgetown: 1 ←
...
```

Values for the `addr:city` tag are a little messy. To fix them, I'll capitalize just the first letter of each word in each city name. A dictionary match should clean up the remainder.

state

```
from osmAudit import key_val_counter

# define filename
atx_filename = 'data/austin_texas.osm'

# key -> addr:state
key_val_counter(atx_filename, 'addr:state')
```

```
----- Count of values for key: addr:state -----
TX: 19311
FL: 1 ←
AL: 1 ←
tx: 1 ←
```

There are a few non-Texas values in this key that need to be filtered out while cleaning and staging the data.

Other Considerations

I have a few additional cleaning steps to integrate into the data preparation function. There are also values for `addr:postcode`, `addr:state`, and `surface` that will be used to remove problematic elements. In addition to this, there are a set of characters that will cause problems when staging this data - any elements with these characters will be removed as well.

Cleaning & Transforming

To prepare the data for my database I need to clean and filter the raw OpenStreetMap data. Then, I need to transform the data from xml format to a tabular format (csv).

Cleaning

First, I wrote a function set for each of the problematic keys I outlined above.

building

For this key, I created a dictionary to correct a few bad values. The `clean_building`⁶ function compares the input value to that dictionary; if the value is contained in the dictionary keys, it's replaced with the dictionary value. Next, the value is checked for spaces, any that are located are replaced with an underscore.

```
building_dict = {
    'Bing': 'yes',
    'Learning_Center/_Day_Care': 'learning_center',
    'sports_centre': 'sports_center'
}

def clean_building(val):
    """Cleans key values for building tag"""
    for key in building_dict.keys():
        if val == key:
            val = building_dict.get(key)
    val = val.replace(' ', '_')
    return val
```

postcode

I created two functions for the `addr:postcode` key:

- The `clean_postcode`⁷ function first takes the input value, splits on semicolon, and drops anything after the semicolon.
 - '12345; 98765' → '12345'
- Next, it drops the last four digits from any values that have the full 9 digit zip code
 - 12345-6789 → 12345

- Then, the `filter_postcode`⁷ function checks a list of valid Austin, TX zip codes.⁸ It returns *False* if that zip code is present on the list (meaning it should not be removed), and *True* if that zip code is not present on the list (meaning it should be removed).

```

atx_postcodes = [
    '78610', '78613', '78617', '78641', '78652', '78653', '78660', '78664',
    '78681', '78701', '78702', '78703', '78704', '78705', '78712', '78717',
    '78719', '78721', '78722', '78723', '78724', '78725', '78726', '78727',
    '78728', '78729', '78730', '78731', '78732', '78733', '78734', '78735',
    '78736', '78737', '78738', '78739', '78741', '78742', '78744', '78745',
    '78746', '78747', '78748', '78749', '78750', '78751', '78752', '78753',
    '78754', '78756', '78757', '78758', '78759'
]

def filter_postcode(val):
    """Filters key values for addr:postcode tag"""

    # run val through postcode cleaning function
    val = clean_postcode(val)

    # set to true if val is not an austin, tx zip code
    if val not in atx_postcodes:
        return True
    else:
        return False

def clean_postcode(val):
    """Cleans key values for addr:postcode tag"""

    # remove multiple zip code entries (e.g. '12345; 98765')
    split_val = val.split('; ', maxsplit=1)
    val = split_val[0]

    # drop last four from full zip codes (e.g. 12345-6789 -> 12345)
    if len(val) == 10:
        val = val[0:5]

    return val

```

surface

I created two functions for the surface key:

- For the `clean_surface`⁹ function, I created a dictionary to correct a few bad values. Then, the input value is compared to that dictionary; if the value is contained in the dictionary keys, it is replaced with the dictionary value.

- The `filter_surface`⁹ function checks a list of values to remove. It returns True if the input value is on that list (meaning it should be removed), and False if the value is not on that list (meaning it should not be removed).

```
surface_dict = {
    'con': 'concrete',
    'large_unattached_stones_through_water': 'stones',
    'Large_unattached_stones_laid_in_the_creek': 'stones',
    'paving_stones:30': 'paving_stones',
    'creekbed_(rock)': 'rock',
    'concrete,_dirt': 'concrete;dirt',
    'dirt/sand': 'dirt;sand',
    'concrete:lanes': 'concrete',
    'concrete:plates': 'concrete'
}

remove_list = [
    'yes', 'CR_127', 'f'
]

def filter_surface(val):
    """Cleans key values for surface tag"""
    if val in remove_list:
        return True
    else:
        return False

def clean_surface(val):
    """Cleans key values for surface tag"""
    for key in surface_dict.keys():
        if val == key:
            val = surface_dict.get(key)
    return val
```

city

The city key required the most cleaning among those I selected, and the function I created, `clean_city`,¹⁰ has multiple steps:

1. First, the function splits any city names that have multiple words.
 - round rock → [round, rock]
2. Then, it capitalizes each of those words by looping through each item in the list created when splitting the value.
 - [round, rock] → [Round, Rock]
3. Next, it puts the city names back together, and retains a space between each word.
 - [Round, Rock] → Round Rock
4. Finally, the value is compared to a dictionary to clean up any lingering incorrect city names.

```

city_dict = {
    'Wells Branch': 'Austin',
    'Barton Creek': 'Austin',
    'Ste 128, Austin': 'Austin',
    'Pepe's Tacos': 'Austin',
    'N Austin': 'Austin',
    'Austin;austin': 'Austin',
    'San Gabriel Village Boulevard': 'Georgetown',
    'Manchaca,' : 'Manchaca',
    'Pflugerville': 'Pflugerville'
}

def clean_city(val):
    """Cleans key values for addr:city tag"""
    split = val.split(' ')
    i = 0
    cap = ''
    while i < len(split):
        x = split[i].capitalize()
        if i == 0:
            cap = x
        else:
            cap = cap + ' ' + x
        i += 1
    val = cap
    for key in city_dict.keys():
        if val == key:
            val = city_dict.get(key)
    return val

```

state

Creating a function just to filter for `addr:state == TX` would have been a textbook example of over-engineering a problem. Instead of creating a function, I just added that filter to the `shape` function outlined below.

Transforming

After the cleaning and filtering functions were developed, I wrote a function, `shape_element`,^{[11](#)} that shapes the node and way elements of the raw xml file, and returns them as a Python dictionary.

Employing that function, the cleaning functions outlined above, and several helper functions provided by Udacity for this project;^{[11](#)} I cleaned, filtered, transformed, and staged the data into csv format to prepare it to be loaded into a sqlite database.

```

import cerberus
import codecs
import csv
import random

```

```

import pprint
import re
import xml.etree.ElementTree as ET
import os

import schema
from osmKeySurface import filter_surface, clean_surface
from osmKeyPostcode import filter_postcode, clean_postcode
from osmKeyCity import clean_city
from osmKeyBuilding import clean_building
from TicToc import TicToc
t = TicToc()

# ===== #
#                               Define Variables                               #
# ===== #

OSM_PATH = 'data/austin_texas.osm'
NODES_PATH = 'data/csv/nodes.csv'
NODE_TAGS_PATH = 'data/csv/nodes_tags.csv'
WAYS_PATH = 'data/csv/ways.csv'
WAY_NODES_PATH = 'data/csv/ways_nodes.csv'
WAY_TAGS_PATH = 'data/csv/ways_tags.csv'

LOWER_COLON = re.compile(r'^([a-z]|_)+:([a-z]|_)+')
PROBLEMCHARS = re.compile(r'[=\/&<>;\'\"?%#$@\,\.\ \t\r\n]')
SCHEMA = schema.schema

NODE_FIELDS = ['id', 'lat', 'lon', 'user', 'uid', 'version', 'changeset', 'timestamp']
NODE_TAGS_FIELDS = ['id', 'key', 'value', 'type']
WAY_FIELDS = ['id', 'user', 'uid', 'version', 'changeset', 'timestamp']
WAY_TAGS_FIELDS = ['id', 'key', 'value', 'type']
WAY_NODES_FIELDS = ['id', 'node_id', 'position']

# ===== #
#                               Shape Function                               #
# ===== #

def shape_element(element):
    """Shape node and way XML elements to Python dictionary"""
    way_attr_fields = WAY_FIELDS
    node_attr_fields = NODE_FIELDS
    problem_chars = PROBLEMCHARS
    default_tag_type = 'regular'

    node_attribs = {}
    way_attribs = {}
    way_nodes = []
    tags = []

    if element.tag == 'node':

```

```

for i in node_attr_fields:
    node_attribs[i] = element.get(i)
for j in element.iter('tag'):
    key = j.get('k')
    val = j.get('v')
    if re.match(problem_chars, key):
        continue
    if key == 'addr:state' and val != 'TX':
        continue
    if key == 'addr:postcode':
        if filter_postcode(val):
            continue
    if key == 'surface':
        if filter_surface(val):
            continue
    if key == 'addr:postcode':
        val = clean_postcode(val)
    elif key == 'addr:city':
        val = clean_city(val)
    elif key == 'building':
        val = clean_building(val)
    elif key == 'surface':
        val = clean_surface(val)
    mat = re.match(LOWER_COLON, key)
    if mat:
        key_split = re.split(':', key, maxsplit=1)
        tags_dict = {
            'id': node_attribs['id'],
            'key': key_split[1],
            'value': val,
            'type': key_split[0]
        }
    else:
        tags_dict = {
            'id': node_attribs['id'],
            'key': key,
            'value': val,
            'type': default_tag_type
        }
    tags.append(tags_dict)
return {
    'node': node_attribs,
    'node_tags': tags
}
elif element.tag == 'way':
    for i in way_attr_fields:
        way_attribs[i] = element.get(i)
    count = 0
    for x in element.iter('nd'):
        way_nodes_dict = {
            'id': way_attribs['id'],

```

```

        'node_id': x.get('ref'),
        'position': count
    }
    count += 1
    way_nodes.append(way_nodes_dict)
for j in element.iter('tag'):
    key = j.get('k')
    val = j.get('v')
    if re.match(problem_chars, key):
        continue
    if key == 'addr:state' and val != 'TX':
        continue
    if key == 'addr:postcode':
        if filter_postcode(val):
            continue
    if key == 'surface':
        if filter_surface(val):
            continue
    if key == 'addr:postcode':
        val = clean_postcode(val)
    elif key == 'addr:city':
        val = clean_city(val)
    elif key == 'building':
        val = clean_building(val)
    elif key == 'surface':
        val = clean_surface(val)
    mat = re.match(LOWER_COLON, key)
    if mat:
        key_split = re.split(':', key, maxsplit=1)
        way_tags_dict = {
            'id': way_attribs['id'],
            'key': key,
            'value': val,
            'type': key_split[0]
        }
    else:
        way_tags_dict = {
            'id': way_attribs['id'],
            'key': key,
            'value': val,
            'type': default_tag_type
        }
    tags.append(way_tags_dict)
return {
    'way': way_attribs,
    'way_nodes': way_nodes,
    'way_tags': tags
}

```

```

# ===== #
# Assistant to the Regional Functions #

```



```

# ===== #

def get_element(osm_file, tags=('node', 'way', 'relation')):
    """Yield element if it is the right type of tag"""
    context = eT.iterparse(osm_file, events=('start', 'end'))
    _, root = next(context)
    for event, elem in context:
        if event == 'end' and elem.tag in tags:
            yield elem
            root.clear()

def validate_element(element, validator, csv_schema=SCHEMA):
    """Raise ValidationError if element does not match schema"""
    if validator.validate(element, csv_schema) is not True:
        field, errors = next(iter(validator.errors.items()))
        message_string = ''\nElement of type '{0}' has the following errors:\n{1}''
        error_string = pprint.pformat(errors)
        raise Exception(message_string.format(field, error_string))

class UnicodeDictWriter(csv.DictWriter, object):
    """Extend csv.DictWriter to handle Unicode input"""
    def writerow(self, row):
        super(UnicodeDictWriter, self).writerow(row)
    def writerows(self, rows):
        for row in rows:
            self.writerow(row)

# ===== #
#                                     Main Function                                     #
# ===== #

def process_map(file_in, validate):
    """Iteratively process each XML element and write to csv(s)"""
    with codecs.open(NODES_PATH, 'w', encoding='utf8') as nodes_file, \
        codecs.open(NODE_TAGS_PATH, 'w', encoding='utf8') as nodes_tags_file, \
        codecs.open(WAYS_PATH, 'w', encoding='utf8') as ways_file, \
        codecs.open(WAY_NODES_PATH, 'w', encoding='utf8') as way_nodes_file, \
        codecs.open(WAY_TAGS_PATH, 'w', encoding='utf8') as way_tags_file:
        nodes_writer = UnicodeDictWriter(nodes_file, NODE_FIELDS)
        node_tags_writer = UnicodeDictWriter(nodes_tags_file, NODE_TAGS_FIELDS)
        ways_writer = UnicodeDictWriter(ways_file, WAY_FIELDS)
        way_nodes_writer = UnicodeDictWriter(way_nodes_file, WAY_NODES_FIELDS)
        way_tags_writer = UnicodeDictWriter(way_tags_file, WAY_TAGS_FIELDS)
        nodes_writer.writeheader()
        node_tags_writer.writeheader()
        ways_writer.writeheader()
        way_nodes_writer.writeheader()
        way_tags_writer.writeheader()
        validator = cerberus.Validator()
        for element in get_element(file_in, tags=('node', 'way')):
            elem = shape_element(element)

```

```

        elem = shape_element(element)
    if elem:
        if validate is True:
            validate_element(elem, validator)
        if element.tag == 'node':
            nodes_writer.writerow(elem['node'])
            node_tags_writer.writerow(elem['node_tags'])
        elif element.tag == 'way':
            ways_writer.writerow(elem['way'])
            way_nodes_writer.writerow(elem['way_nodes'])
            way_tags_writer.writerow(elem['way_tags'])

# ===== #
#                                     Execute                                     #
# ===== #

if __name__ == '__main__':
    py = os.path.basename(__file__)
    print('\nExecuting ' + py + '....')
    t.tic()
    process_map(OSM_PATH, validate=False)
    t.toc()

```

Problems Encountered

I encountered several problems while working with the Austin OpenStreetMap data. Chief among them was file size. I did not anticipate how resource intensive working with a dataset of this size would be. If I were to do this project again I would select a smaller map to work with. As you can see, some of the files used are quite large.^{[12](#)}

```

import os
import pandas as pd
from TicToc import TicToc
t = TicToc()

# ===== #
#                                     Define Variables                                     #
# ===== #

root_data = '/Users/ajp/dsProjects/workspace/osmAustin/data/'
root_csv = root_data + 'csv/'

osm = 'austin_texas.osm'
sample = 'sample_atx.osm'
nodes_tags = 'nodes_tags.csv'
nodes = 'nodes.csv'
ways_nodes = 'ways_nodes.csv'
ways_tags = 'ways_tags.csv'
ways = 'ways.csv'

```

```

path_osm = root_data + osm
path_sample = root_data + sample
path_nodes_tags = root_csv + nodes_tags
path_nodes = root_csv + nodes
path_ways_nodes = root_csv + ways_nodes
path_ways_tags = root_csv + ways_tags
path_ways = root_csv + ways

# ===== #
#                                     Size Function                                     #
# ===== #

def get_size(filepath):
    """Get file size and return string with appropriate unit"""
    size_bytes = os.path.getsize(filepath)
    if size_bytes < 1024:
        size_bytes = round(size_bytes, 2)
        size = f'{size_bytes} B'
    else:
        size_kilobytes = size_bytes / 1024
        if size_kilobytes < 1024:
            size_kilobytes = round(size_kilobytes, 2)
            size = f'{size_kilobytes} KB'
        else:
            size_megabytes = size_kilobytes / 1024
            if size_megabytes < 1024:
                size_megabytes = round(size_megabytes, 2)
                size = f'{size_megabytes} MB'
            else:
                size_gigabytes = size_megabytes / 1024
                if size_gigabytes < 1024:
                    size_gigabytes = round(size_gigabytes, 2)
                    size = f'{size_gigabytes} GB'
                else:
                    size = "Wow, that's huge."
    return size

# ===== #
#                                     Execute                                     #
# ===== #

if __name__ == '__main__':
    py = os.path.basename(__file__)
    print('\nExecuting ' + py + '....')
    t.tic()

    osm_size = get_size(path_osm)
    sample_size = get_size(path_sample)
    nodes_tags_size = get_size(path_nodes_tags)
    nodes_size = get_size(path_nodes)

```

```

nodes_size = get_size(path_nodes)
ways_nodes_size = get_size(path_ways_nodes)
ways_tags_size = get_size(path_ways_tags)
ways_size = get_size(path_ways)

names = [osm, sample, nodes_tags, nodes, ways_nodes, ways_tags, ways]
sizes = [osm_size, sample_size, nodes_tags_size, nodes_size, ways_nodes_size, ways_tags_size, ways_size]

sizes_dict = {
    'name': names,
    'size': sizes
}

sizes_df = pd.DataFrame(data=sizes_dict)
print(sizes_df)
t.toc()

```

	name	size
0	austin_texas.osm	1.66 GB
1	sample_atx.osm	40.06 MB
2	nodes_tags.csv	13.8 MB
3	nodes.csv	696.06 MB
4	ways_nodes.csv	203.81 MB
5	ways_tags.csv	85.79 MB
6	ways.csv	56.83 MB

To work through this problem, I created a sample file that only contains elements with ways tags for `addr:state=TX`. This step significantly sped up testing and development, reducing the working file size from 1.79 GB to 42 MB. I did this with the java-based `osmosis` tool used earlier in this document.^{[3](#)}

```

# get sample dataset for testing/development
osmosis --rx file=austin_texas.osm --tf accept-ways addr:state=TX --un --wx sample_atx.osm

```

I also had a little trouble finding data to clean. Many of the keys' values were already very clean. I suspect that since Austin is a tech city other students and hobbyists like myself have done similar projects and cleaned the OpenStreetMap data for this metropolitan area.

Sqlite Database

After the data was cleaned, I created a sqlite database.^{[4](#)} Then, I created a schema in that database to match the schema of my csv files. After that, I loaded the data into their respective tables.

```
create table nodes
(
  id          integer not null
              constraint nodes_pk
              primary key,
  lat         real,
  lon         real,
  user        text,
  uid         integer,
  version     integer,
  changeset   integer,
  timestamp   text
);
```

```
create table nodes_tags
(
  id          integer
              references nodes,
  key         text,
  value       text,
  type        text
);
```

```
create table ways
(
  id          integer not null
              constraint ways_pk
              primary key,
  user        text,
  uid         integer,
  version     text,
  changeset   integer,
  timestamp   text
);
```

```
create table ways_nodes
(
  id          integer not null
              references ways,
  node_id     integer not null
              references nodes,
  position    integer not null
);
```

```
create table ways_tags
(
  id      integer not null
        references ways,
  key     text     not null,
  value   text     not null,
  type    text
);
```

Overview of the Data

SQL Stats

Before digging into the dataset, I took a look at some database stats to get an idea of how much data I'd be working with. To generate those stats, I used another command-line utility program called `sqlite3_analyzer`.^{[15](#)}

```
sqlite3_analyzer --stats osmDB.sqlite
```

Then, I loaded the stats into a table in my database.^{[16](#)}

The insert statements for each set of statistics are extremely long, so I'll leave them out of this document. However, they can be found in this repo for reference.^{[16](#)}

```
BEGIN;
CREATE TABLE stats(
  name      STRING,          /* Name of table or index */
  path      INTEGER,         /* Path to page from root */
  pageno    INTEGER,         /* Page number */
  pagetype  STRING,          /* 'internal', 'leaf' or 'overflow' */
  ncell     INTEGER,         /* Cells on page (0 for overflow) */
  payload   INTEGER,         /* Bytes of payload on this page */
  unused    INTEGER,         /* Bytes of unused space on this page */
  mx_payload INTEGER,        /* Largest payload size of all cells */
  pgoffset  INTEGER,         /* Offset of page in file */
  pgsize    INTEGER          /* Size of the page */
);
COMMIT;
```

After the stats table was created, I wrote a simple query to check table sizes.^{[17](#)}

```

SELECT name AS table_name
      , CASE
          WHEN bytes < 1024 THEN (bytes || ' B')
          WHEN kilobytes < 1024 THEN ROUND(kilobytes, 2) || ' KB'
          ELSE ROUND(megabytes, 2) || ' MB' END AS table_size
FROM (
    SELECT name
          , SUM(payload) as bytes
          , CAST(SUM(payload) AS FLOAT) / 1024 AS kilobytes
          , (CAST(SUM(payload) AS FLOAT) / 1024) / 1024 AS megabytes
    FROM stats
    GROUP BY name
)
ORDER BY bytes DESC;

```

	table_name	table_size
1	nodes	528.53 MB
2	ways_nodes	123.58 MB
3	ways_tags	73.9 MB
4	ways	42.98 MB
5	nodes_tags	12.23 MB
6	sqlite_schema	2.33 KB

Analysis

Now that the stats were collected I could start analyzing the clean data. As I was analyzing this dataset, I wrote several queries¹⁸ to answer some investigative questions:

How many people have contributed to the Austin OpenStreetMap project?

```

SELECT COUNT(DISTINCT uid) AS contributors
FROM (
    SELECT uid FROM nodes
    UNION ALL
    SELECT uid FROM ways
);

```

	contributors
1	2,973

Who are the top contributors?

```
SELECT uid,
       user,
       COUNT(*) AS contributions
FROM (
    SELECT uid, user FROM nodes
    UNION ALL
    SELECT uid, user FROM ways
)
GROUP BY uid, user
ORDER BY contributions DESC
LIMIT 10;
```

	uid	user	contributions
1	3369502	patisilva_atxbldings	2,638,615
2	3370181	ccjmartin_atxbldings	1,257,245
3	3405475	ccjmartin_atxbldings	920,175
4	3341346	wilsaj_atxbldings	349,180
5	3341321	jseppi_atxbldings	284,518
6	147510	woodpeck_fixbot	179,918
7	3367383	kkt_atxbldings	155,199
8	3409435	lyzidiamond_atxbldings	150,603
9	5446055	torapa	131,329
10	12056179	Milroy1812	124,175

How many total contributions are made each year?

```
SELECT year,
       COUNT(*) AS contributions
FROM (
    SELECT SUBSTR(timestamp, 1, 4) AS year FROM nodes
    UNION ALL
    SELECT SUBSTR(timestamp, 1, 4) AS year FROM ways
)
GROUP BY year
ORDER BY year;
```

	year	contributions
--	------	---------------

1	2007	2,111
2	2008	18,020
3	2009	223,961
4	2010	20,007
5	2011	44,811
6	2012	112,824
7	2013	62,621
8	2014	97,015
9	2015	5,958,603
10	2016	73,128
11	2017	79,173
12	2018	227,833
13	2019	581,022
14	2020	534,813
15	2021	754,611

Which years had the most contributions?

```
SELECT year,
       COUNT(*) AS contributions
FROM (
    SELECT SUBSTR(timestamp, 1, 4) AS year FROM nodes
    UNION ALL
    SELECT SUBSTR(timestamp, 1, 4) AS year FROM ways
)
GROUP BY year
ORDER BY contributions DESC
LIMIT 3;
```

	year	contributions
1	2015	5,958,603
2	2021	754,611
3	2019	581,022

During which months are the contributors most active?

```
SELECT CASE
    WHEN mon = '01' THEN 'January'
    WHEN mon = '02' THEN 'February'
    WHEN mon = '03' THEN 'March'
    WHEN mon = '04' THEN 'April'
    WHEN mon = '05' THEN 'May'
    WHEN mon = '06' THEN 'June'
    WHEN mon = '07' THEN 'July'
    WHEN mon = '08' THEN 'August'
    WHEN mon = '09' THEN 'September'
    WHEN mon = '10' THEN 'October'
    WHEN mon = '11' THEN 'November'
    WHEN mon = '12' THEN 'December'
END AS month,
COUNT(*) AS contributions
FROM (
    SELECT SUBSTR(timestamp, 6, 2) AS mon FROM nodes
    UNION ALL
    SELECT SUBSTR(timestamp, 6, 2) AS mon FROM ways
)
GROUP BY month
ORDER BY mon
LIMIT 3;
```

	month	contributions
1	May	153,204
2	September	211,860
3	April	215,569

What are the average monthly and yearly contributions?

```

SELECT SUM(contributions) / COUNT(DISTINCT year) AS avg_yearly,
       SUM(contributions) / COUNT(DISTINCT month) AS avg_monthly
FROM (
  SELECT year,
         month,
         COUNT(*) AS contributions
  FROM (
    SELECT SUBSTR(timestamp, 1, 4) AS year,
           SUBSTR(timestamp, 1, 7) AS month
    FROM nodes
    UNION ALL
    SELECT SUBSTR(timestamp, 1, 4) AS year,
           SUBSTR(timestamp, 1, 7) AS month
    FROM ways
  )
  GROUP BY year, month
);

```

	avg_yearly	avg_monthly
1	586,036	53,929

How many total ways tags are in this dataset, and what are the most common tags?

```

-- count the ways
SELECT COUNT(*) AS count_ways
FROM ways;

-- common way tags
SELECT key,
       COUNT(*) AS count_ways
FROM ways_tags
GROUP BY key
HAVING count_ways > 100000
ORDER BY count_ways DESC;

```

	key	count_ways
-	Total	858,496
1	building	620,762
2	height	438,569
3	addr:street	261,635
4	addr:housenumber	260,847
5	highway	189,824

	key	count_ways
--	-----	------------

How many total nodes tags are in this dataset, and what are the most common tags?

```
-- count the nodes
SELECT COUNT(*) AS count_nodes
FROM nodes;

-- common node tags
SELECT key,
       COUNT(*) AS count_nodes
FROM nodes_tags
GROUP BY key
HAVING count_nodes > 10000
ORDER BY count_nodes DESC;
```

	key	count_nodes
-	Total	7,932,057
1	street	83,171
2	houenumber	83,135
3	postcode	62,267
4	highway	26,828
5	barrier	16,292
6	power	13,811

What are the most common amenities in Austin, TX?

```
SELECT value,
       COUNT(*) AS count
FROM nodes_tags
WHERE key='amenity'
GROUP BY value
HAVING count >= 200
ORDER BY count DESC;
```

	value	count
1	restaurant	916
2	bench	887
3	waste_basket	791

4	fast_food	427
5	loading_dock	316
6	parking_entrance	257
7	place_of_worship	231
8	bicycle_parking	219

What kind of restaurants are popular in Austin?

```
SELECT value,
        COUNT(*) AS count_restaurants
FROM nodes_tags
WHERE key = 'cuisine'
GROUP BY value
ORDER BY count_restaurants DESC
LIMIT 10;
```

	value	count_restaurants
1	sandwich	116
2	pizza	115
3	mexican	110
4	coffee_shop	77
5	burger	66
6	chinese	40
7	american	38
8	indian	29
9	thai	25
10	italian	24

Other Ideas About the Dataset

I think one of the best things OpenStreetMap could do to improve their data would be to develop a robust, automated cleaning process. These scripts or bots could automatically make corrections of specific errors.

It looks like there is a bot (woodpeck_fixbot) modifying elements in the Austin dataset, but the scope of that project must be relatively narrow because I still found simple corrections to make in the data.

The kind of bot I'm imagining is more extensive, one example would be Wall-E,^{[19](#)} but it is currently only operational in Germany and Austria. Perhaps OSM is worried that an automated correction bot for a map as large as the United States could cause problems if it was making inaccurate 'corrections'.

Those hurdles could be overcome through proper testing. Specifically, by testing features on a very small area, and slowly widening the bot's scope before unleashing it on the entire map.

Conclusion

These data went on a long journey before landing in a clean sqlite database.

- I converted pbf file to osm format using osmosis, a java-based command line tool.
- Then I investigated the raw data in a python notebook.
- Next, I cleaned, filtered, and transformed the data using a handful of python scripts and functions.
- Finally, I loaded the data into a sqlite database and analyzed it with SQL.

There is additional work that could be done on the OpenStreetMap data for Austin, TX. Also, data issues are likely to be a constant problem for OSM until someone implements a more widespread automated cleaning process.

```

  ----
 /  _ )-----  _ _ _ _ _ / _ _ / _ _ _ _ _ / /
 /  _ / _ _ / _ _ \ / _ _ \ / _ _ // _ \ / / _ _ \ / _ _ \ / _ _ \
 / / _ / / / / / _ / / / / / / / _ / / / / _ / / _ / / _ / /
 / _ _ _ / / \ _ _ _ / / / / _ _ | _ _ / / \ _ _ _ / \ _ _ , \ _ _ , /
```